

Cloud Computing
PA 1
Benchmarking
Design Document and Performance

Name: Sharanheer Choudhari (A20398615)

Design Document

This assignment aims at Benchmarking different parts of the computer system using the virtual cluster running on the chameleon testbed.

CPU Benchmarking

- For CPU Benchmarking, I used strong scaling experiment and design a source code in C language.
- This program will measure the processor speed, in terms of quarter precision, half precision, single precision and double precision floating point operations.(GOPS/Sec)
- This program has different function each consisting of 100 instruction which include addition, subtraction, multiplication operation on char, short, integer and floating point numbers which will run for 1 Trillion times to measure the speed.
- Main aim of this is to utilize the complete CPU cycle by performing diverse types of instructions and measure the speed.
- There is a performance evaluation which evaluates OPS executed per second using different levels of concurrency (1 thread, 2 threads, 4 threads).
- Ran the Linpack benchmark (<http://en.wikipedia.org/wiki/LINPACK>) and compare the performance achieved using double precision floating point.

Test Results: Performed experiment on Prometheus Cluster for CPU benchmark.

Workload	Concurrency	MyCPUBench Measured Ops/Sec (GigaOPS)	HPL Measured Ops/Sec (GigaOPS)	Theoretical Ops/Sec (GigaOPS)	MyCPUBench Efficiency (%)	HPL Efficiency (%)
QP	1	1.12	N/A	537.6	0.002	N/A
QP	2	2.15	N/A	537.6	0.0039	N/A
QP	4	2.17	N/A	537.6	0.00403	N/A
HP	1	1.16	N/A	537.6	0.0022	N/A
HP	2	2.24	N/A	537.6	0.0042	N/A
HP	4	2.13	N/A	537.6	0.0039	N/A
SP	1	2.11	N/A	537.6	0.0039	N/A
SP	2	3.97	N/A	537.6	0.0074	N/A
SP	4	4.08	N/A	537.6	0.0076	N/A
DP	1	2.06	421.79	537.6	0.0038	78.45
DP	2	4.12	421.79	537.6	0.0076	78.45
DP	4	3.66	421.79	537.6	0.0068	78.45

Theoretical Performance

Based on the hardware specification of the Prometheus Cluster.

FLOPS = Number of Cores * CLOCK SPEED * IPC * Number of CPU's

$$= 8 * 2.1 * 10^6 * 4 * 8 = 537.6 \text{ GFLOPS}$$

Linpack Benchmark:

Performed experiment on Hyperion Cluster for CPU benchmark.

I have ran the LINPACK Benchmark on the system to compare the CPU performance and from the below it can seen that the maximal value is obtained at 429.27 GFLOPS which is 79.8% matching with the theoretical performance.

```
schoudhari@hyperionides:~/linpack_11.1.2/benchmarks/linpack$ ./xlinpack_xeon64
Input data or print help ? Type [data]/help :
1000000000
Number of equations to solve (problem size): 10000
Leading dimension of array: 10000
Number of trials to run: 10
Data alignment value (in Kbytes): 10000
Current date/time: Tue Mar 27 21:59:04 2018

CPU frequency:      3.090 GHz
Number of CPUs: 16
Number of cores: 16
Number of threads: 16

Parameters are set to:

Number of tests: 1
Number of equations to solve (problem size) : 10000
Leading dimension of array                  : 10000
Number of trials to run                    : 10
Data alignment value (in Kbytes)           : 10000

Maximum memory requested that can be used=810440000, at the size=10000

===== Timing linear equation system solver =====

Size  LDA   Align. Time(s)   GFlops   Residual   Residual(norm) Check
10000 10000 10000   1.607    414.9071 7.750056e-11 2.732748e-02 pass
10000 10000 10000   1.570    424.6473 7.750056e-11 2.732748e-02 pass
10000 10000 10000   1.563    426.7138 7.750056e-11 2.732748e-02 pass
10000 10000 10000   1.589    419.5710 7.750056e-11 2.732748e-02 pass
10000 10000 10000   1.596    417.8517 7.750056e-11 2.732748e-02 pass
10000 10000 10000   1.615    412.9670 7.750056e-11 2.732748e-02 pass
10000 10000 10000   1.553    429.2731 7.750056e-11 2.732748e-02 pass
10000 10000 10000   1.557    428.1991 7.750056e-11 2.732748e-02 pass
10000 10000 10000   1.602    416.3675 7.750056e-11 2.732748e-02 pass
10000 10000 10000   1.560    427.4160 7.750056e-11 2.732748e-02 pass

Performance Summary (GFlops)

Size  LDA   Align. Average Maximal
10000 10000 10000   421.7914 429.2731

Residual checks PASSED

End of tests

schoudhari@hyperionides:~/linpack_11.1.2/benchmarks/linpack$
```

Memory Benchmark:

- For Memory Benchmarking, I used strong scaling experiment and design a source code in C language.
- This program will measure the memory speed of host which will include parameters like read + write operations, sequential read-write access, random read-write access.
- This experiment used different block sizes (1B, 1KB, 1MB, 10MB), and varying the concurrency (1thread, 2 threads and 4 threads).
- We allocate 1 GB of memory (x100) to perform read + write operations in both sequential and random access.
- The program also measures the throughput (Megabytes per second, MB/sec) and latency(microseconds, us). 1B block case used to measure latency, and the 1KB, 1MB, and 10MB cases used to measure throughput.
- Computed the theoretical memory bandwidth of the cluster memory, based on the type of memory and the speed.
- Ran the PMBW benchmark (preinstalled on the cluster) and compared both the performance to see weather benchmarking is running properly.

Test Output: Performed experiment on Prometheus Cluster for memory benchmark.

Throughput Reading

Work-load	Con-curren-cy	Block Size	MyRAMBench Measured Throughput (GB/sec)	pmbw Measured Throughput (GB/sec)	Theoretical Throughput (GB/sec)	MyRAMBench Efficiency (%)	pmbw Efficiency (%)
RWS	1	1KB	3.93	9.8	33.23	11.82	26.72
RWS	1	1MB	3.56	16.86	29.45	12.08	57.25
RWS	1	10MB	4.33	16.88	31.12	13.91	54.24
RWS	2	1KB	6.04	30.7	52.50	11.504	58.47
RWS	2	1MB	4.98	32.17	40.12	12.41	80.18
RWS	2	10MB	5.224	33.09	41.23	12.67	80.25
RWS	4	1KB	6.26	35.6	54.43	11.5	65.40
RWS	4	1MB	5.77	36.46	43.23	13.34	84.33
RWS	4	10MB	6.08	31.89	45.34	13.40	70.33
RWR	1	1KB	2.49	6.7	22.19	11.22	30.19
RWR	1	1MB	3.61	14.34	25.89	13.94	55.38
RWR	1	10MB	3.68	15.23	24.56	14.98	62.01
RWR	2	1KB	1.33	23.45	19.34	5.44	95.98
RWR	2	1MB	5.47	24.43	39.74	13.76	65.47
RWR	2	10MB	5.69	25.43	38.56	14.75	65.94
RWR	4	1KB	1.91	27.89	16.23	6.7	97.96
RWR	4	1MB	5.90	28.19	48.33	12.20	58.32
RWR	4	10MB	6.23	28.79	46.97	13.27	61.29

Latency Reading:

Work-load	Con-curren-cy	Blo-ck Size	MyRAMBenc-h Measured Latency (us)	pmbw Measure-d Latency	Theoretic-al Latency (us)	MyRAMBenc-h Efficiency (%)	pmbw Efficien-cy (%)
RWS	1	1B	0.004	0.027	0.031	12.90	87.09
RWS	2	1B	0.0026	0.009	0.014	18.57	64.28
RWS	4	1B	0.002251	0.008	0.012	18.75	66.66
RWR	1	1B	0.13	0.583	0.783	16.60	74.45
RWR	2	1B	0.35	0.562	0.89	39.32	63.15
RWR	4	1B	0.399	0.619	0.93	42.90	66.55

Theoretical Performance:

Throughput = (Data Transfer/Clock) * Clock Frequency * Memory bus width * Number of interfaces

PMBW BENCHMARK:

Below is the screenshot of PMBW benchmarking showing the results of Memory Speeds. The result below shows the memory speed having a total of 23.54GB/s.

```
schoudhari@compute-3:~$ pmbw -p 2 -s 1024 -S 1024
Running benchmarks with at least 2 threads.
Running benchmarks with array size at least 1024.
Running benchmarks with array size up to 1024.
CPUID: mmx sse avx
Detected 3951 MiB physical RAM and 2 CPUs.

Allocating 2048 MiB for testing.
Running nthreads=2 factor=1073741824 areastestsize=1280 repeats=1677722 testvol=2147484160 testaccess=268435520
run time = 0.101292 -> rerunning test with repeat factor=15900620294
Running nthreads=2 factor=15900620294 areastestsize=1280 repeats=24844720 testvol=31801241600 testaccess=3975155200
run time = 1.35052 -> next test with repeat factor=17660564596
RESULT datetime=2018-03-28 03:23:24 host=compute-3 version=0.6.2 funcname=ScanWrite64PtrSimpleLoop nthreads=2
areastestsize=1024 thrsizetestsize=1280 repeats=24844720 testvol=31801241600 testaccess=3975155200
time=1.3505191789881791919 bandwidth=23547419462.658626556 rate=3.3973998775901356956e-10
Skipping ScanWrite64PtrSimpleLoop test with 2048 maximum array size due to -S 1024.
Skipping ScanWrite64PtrSimpleLoop test with 3072 maximum array size due to -S 1024.
Skipping ScanWrite64PtrSimpleLoop test with 4096 maximum array size due to -S 1024.
Skipping ScanWrite64PtrSimpleLoop test with 6144 maximum array size due to -S 1024.
Skipping ScanWrite64PtrSimpleLoop test with 8192 maximum array size due to -S 1024.
Skipping ScanWrite64PtrSimpleLoop test with 10240 maximum array size due to -S 1024.
```

```
schoudhari@compute-3:~$ pmbw -p 1 -s 1024 -S 1024
Running benchmarks with at least 1 threads.
Running benchmarks with array size at least 1024.
Running benchmarks with array size up to 1024.
CPUID: mmx sse avx
Detected 3951 MiB physical RAM and 2 CPUs.

Allocating 2048 MiB for testing.
Running nthreads=1 factor=1073741824 areastestsize=1152 repeats=932068 testvol=1073742336 testaccess=134217792
run time = 0.065818 -> rerunning test with repeat factor=24470723674
Running nthreads=1 factor=24470723674 areastestsize=1152 repeats=21241948 testvol=24470724096 testaccess=3058840512
run time = 1.57442 -> next test with repeat factor=23314089693
RESULT datetime=2018-03-28 03:22:19 host=compute-3 version=0.6.2 funcname=ScanWrite64PtrSimpleLoop nthreads=1
areastestsize=1024 thrsizetestsize=1152 repeats=21241948 testvol=24470724096 testaccess=3058840512
time=1.5744164420175366104 bandwidth=15542726462.283370972 rate=5.1471020991157075161e-10
Skipping ScanWrite64PtrSimpleLoop test with 2048 maximum array size due to -S 1024.
Skipping ScanWrite64PtrSimpleLoop test with 3072 maximum array size due to -S 1024.
Skipping ScanWrite64PtrSimpleLoop test with 4096 maximum array size due to -S 1024.
Skipping ScanWrite64PtrSimpleLoop test with 6144 maximum array size due to -S 1024.
Skipping ScanWrite64PtrSimpleLoop test with 8192 maximum array size due to -S 1024.
Skipping ScanWrite64PtrSimpleLoop test with 10240 maximum array size due to -S 1024.
```

Disk Benchmarking

- For Disk Benchmarking, I used strong scaling experiment and design a source code in C language.
- This program will measure the disk speed of host which will include parameters like read + write operations, sequential read-write access, random read-write access.
- This experiment used different block sizes (1KB, 1MB, 10MB, 100MB), and varying the concurrency (1thread, 2 threads, 4 threads) for throughput and (thread 1, thread 2, thread 4, thread 8, thread 16, thread 32, thread 64, thread 128) for Latency.
- It implements 4 methods, sequential read and write, random read and write.
- Allocated a large piece file of 10GB, and perform read and write operations on either sequential or random access within this 10GB file.
- Ran the IOZone benchmark (<http://www.iozone.org/>) and compared the performance achieved with theoretical performance.

Test Output: Performed experiment on Hyperion Cluster for disk benchmark.

Throughput Reading

Work-load	Concurren-cy	Block Size	MyDiskBenc h Measured Throughput (MB/sec)	IOZone Measured Throughp ut (MB/	Theoretical Throughput (MB/sec)	MyDiskBe nch Efficiency (%)	IOZone Efficien cy (%)
RS	1	1MB	137.18	341.9	593.37	23.19	57.62
RS	1	10MB	129.115	492.5	685.32	18.84	71.86
RS	1	100M	217.318	783.7	876.31	24.79	89.43
RS	2	1MB	150.07	492.179	657.04	22.84	74.90
RS	2	10MB	132.018	605.7	789.60	16.71	76.71
RS	2	100M	226.553	754.9	891.01	25.42	84.72
RS	4	1MB	162.86	847.23	1039.54	15.66	81.50
RS	4	10MB	197.18	646.6	998.42	19.74	64.76
RS	4	100M	461.906	817.03	1612.3	28.64	50.67
WS	1	1MB	43.91	196.6	302.49	14.51	64.99
WS	1	10MB	60.22	213.4	450.65	23.36	47.35
WS	1	100M	116.74	281.5	527.15	22.15	53.40
WS	2	1MB	44.04	289.86	350.3	12.57	82.74
WS	2	10MB	59.365	223.3	409.71	14.48	54.50
WS	2	100M	118.36	580.65	711.13	16.64	81.65
WS	4	1MB	43.67	355.9	432.21	10.10	82.34
WS	4	10MB	59.39	269.6	509.91	11.65	52.87
WS	4	100M	122.84	511.16	761.38	16.14	67.14
RR	1	1MB	558.83	247.9	357.54	30.57	69.33
RR	1	10MB	292.43	379.05	504.67	57.94	75.11
RR	1	100M	333.73	587.04	891.23	37.45	65.86
RR	2	1MB	181.363	422.7	534.32	33.94	79.10
RR	2	10MB	170.04	608.6	823.41	20.65	73.91
RR	2	100M	220.73	727.15	998.03	22.11	72.85
RR	4	1MB	172.79	778.9	937.40	18.43	83.09
RR	4	10MB	177.84	854.43	1120.32	15.87	76.26

RR	4	100M	172.85	920.63	1398.30	12.36	65.84
WR	1	1MB	51.15	267.4	398.09	12.84	67.17
WR	1	10MB	168.71	273.6	546.26	30.88	50.08
WR	1	100M	192.311	790.08	901.25	21.34	87.66
WR	2	1MB	48.140	292.3	402	11.97	72.71
WR	2	10MB	157.61	444.5	679.30	23.20	65.43
WR	2	100M	153.392	833.8	1023.43	14.98	81.47
WR	4	1MB	46.598	294.5	309.21	15.07	95.24
WR	4	10MB	141.42	417.23	661.23	21.38	63.09
WR	4	100M	132.754	510.21	729.41	18.20	69.94

Latency Reading

Work-load	Con-curren-cy	Blo-ck Size	MyDiskBenc-h Measured Latency (ms)	IOZone Measure-d Latency	Theoretic-al Latency (ms)	MyDiskBenc-h Efficiency (%)	IOZone Efficien-cy (%)
RR	1	1KB	0.456	0.956	1.9	24	50.31
RR	2	1KB	0.165	0.665	0.83	19.87	80.12
RR	4	1KB	0.0667	0.367	0.55	12.12	66.72
RR	8	1KB	0.026	0.096	0.174	14.97	55.17
RR	16	1KB	0.001251	0.0091	0.0121	10.33	75.21
RR	32	1KB	0.00104	0.00804	0.0083	12.53	96.86
RR	64	1KB	0.0142	0.042	0.0787	18.04	53.36
RR	128	1KB	0.01176	0.04176	0.054	21.7	77.33
WR	1	1KB	0.5913	1.913	2.4	24.63	79.70
WR	2	1KB	0.630	1.63	2.3	27.39	70.86
WR	4	1KB	0.6705	1.4605	1.65	40.63	88.51
WR	8	1KB	0.6634	0.934	1.15	57.68	81.21
WR	16	1KB	0.662	1.462	2.23	29.68	65.56
WR	32	1KB	0.602	1.602	2.71	22.21	59.11
WR	64	1KB	0.577	1.577	1.69	34.14	93.31
WR	128	1KB	0.64	1.64	2.12	30.18	77.35

IOPS Reading

Work-load	Con-curren-cy	Blo-ck Size	MyDiskBenc-h Measured IOPS	IOZone Measure-d IOPS	Theoretic-al IOPS	MyDiskBenc-h Efficiency (%)	IOZone Efficien-cy (%)
RR	1	1KB	2192.89	7192.89	8112	27.03	88.66
RR	2	1KB	6031.45	9031.45	16215	37.19	55.69
RR	4	1KB	14975.083	24975.083	35016.32	42.76	71.32
RR	8	1KB	38112.04	48112.04	73084	52.14	65.83
RR	16	1KB	799469.37	899469.37	921554.93	86.75	97.61
RR	32	1KB	958307.215	1258307.215	1794210.10	53.41	70.13
RR	64	1KB	69971.151	149971.151	201923.2	34.65	74.27
RR	128	1KB	84987.869	284987.869	326751.2	26.01	87.21
WR	1	1KB	1690.97	6690.97	8262.1	20.46	80.98
WR	2	1KB	1586.23	7586.23	9012.5	17.60	84.17
WR	4	1KB	1491.38	8491.38	9243.1	16.13	91.86
WR	8	1KB	1507.219	7507.219	8200.1	18.38	91.55
WR	16	1KB	1510.23	7510.23	7832	19.28	95.89
WR	32	1KB	1660.2	7660.2	9271	17.907	82.62
WR	64	1KB	1732.4	6732.4	9769	17.73	68.91
WR	128	1KB	1539.35	5539.35	7865	19.57	70.43

IoZone Benchmarking:

The following screenshot shows the performance metrics taken from IoZone benchmarking ran on the system.

The speeds are measured in terms of Kilobytes. The image shows the speeds for pread & pwrite which runs for block size of 10MB and file size 1GB.

```

schoudhart@bluecompute-1:~$ iozone -l 9 -l 10 -t 4 -s 1G -r 10m
Iozone: Performance Test of File I/O
Version $Revision: 3.471 $
Compiled for 64 bit mode.
Build: linux-AMD64

Contributors:William Norcott, Don Capps, Isom Crawford, Kirby Collin
Al Slater, Scott Rhine, Mike Wisner, Ken Goss
Steve Landherr, Brad Smith, Mark Kelly, Dr. Alain CYR,
Randy Dunlap, Mark Montague, Dan Million, Gavin Brebner
Jean-Marc Zucconi, Jeff Blomberg, Benny Halevy, Dave Bo
Erik Habbinga, Kris Strecker, Walter Wong, Joshua Root,
Fabrice Bacchella, Zhenghua Xue, Qin Li, Darren Sawyer,
Vangel Bojaxhi, Ben England, Vikentsi Lapa,
Alexey Skidanov.

Run began: Wed Mar 28 16:33:11 2018

File size set to 1048576 kB
Record Size 10240 kB
Command line used: iozone -l 9 -l 10 -t 4 -s 1G -r 10m
Output is in kbytes/sec
Time Resolution = 0.000001 seconds.
Processor cache size set to 1024 kbytes.
Processor cache line size set to 32 bytes.
File stride size set to 17 * record size.
Throughput test with 4 processes
Each process writes a 1048576 kbyte file in 10240 kbyte records

Children see throughput for 4 pwrite writers      = 278405.30 kB/sec
Parent sees throughput for 4 pwrite writers      = 261891.60 kB/sec
Min throughput per process                      = 66924.08 kB/sec
Max throughput per process                      = 72088.92 kB/sec
Avg throughput per process                      = 69601.32 kB/sec
Min xfer                                        = 972800.00 kB

Children see throughput for 4 pread readers      = 186435.58 kB/sec
Parent sees throughput for 4 pread readers      = 183999.35 kB/sec
Min throughput per process                      = 46000.48 kB/sec
Max throughput per process                      = 47568.16 kB/sec
Avg throughput per process                      = 46608.89 kB/sec
Min xfer                                        = 1024000.00 kB

iozone test complete.

```


Network Benchmarking

- For Network Benchmarking, I used strong scaling experiment and design a source code in C language.
- It will measure the network speed over the between 2 processes on the different nodes.
- It includes the TCP protocol stack, UDP, varying packet/buffer size (1KB, 32KB, 1Byte), and varying the concurrency (1thread, 2 threads, 4 threads, 8 threads) on both server and client side.
- It will measure the throughput (Megabits per second, Mb/sec) and latency (ms).
- We used fixed amount of data (1GB x100 times) to transfer using socket programming and multithreading.
- Ran the IPerf benchmark (<http://en.wikipedia.org/wiki/Iperf>) and compared achieved throughput performance over the cluster network with theoretical values.

Test Reading Performed experiment on Hyperion Cluster for network benchmark.

Throughput Reading

Prot o- col	Con- curren cy	Blo ck Size	MyNETBenc h Measured Throughput (Mb/sec)	iperf Measured Throughp ut (Mb/ sec)	Theoretical Throughput (Mb/sec)	MyNETBe nch Efficiency (%)	iperf Efficien cy (%)
TCP	1	1KB	172.041	551	593.37	28.99	92.85
TCP	1	32KB	171.956	623.4	685.32	25.09	90.96
TCP	2	1KB	413.15	651	976.31	42.31	66.67
TCP	2	32KB	207.835	721.3	657.04	27.45	95.27
TCP	4	1KB	386.49	739.3	989.6	39.05	74.70
TCP	4	32KB	389.328	821.67	891.01	43.69	92.21
TCP	8	1KB	521.817	904.3	1239.54	42.09	72.95
TCP	8	32KB	531.48	839	998.42	53.23	84.03
UDP	1	1KB	157.824	720.1	912.3	17.29	78.93
UDP	1	32KB	603.232	1387	1602.49	37.64	86.55
UDP	2	1KB	184.512	834	950.65	19.4	87.73
UDP	2	32KB	744.18	1550	1927.15	38.61	80.42
UDP	4	1KB	177.77	893	950.3	18.7	93.97
UDP	4	32KB	712.812	1250	1309.71	54.42	95.44
UDP	8	1KB	156.612	851	911.13	17.18	93.4
UDP	8	32KB	987	1943	2332.21	42.32	83.21

Latency Reading

Prot o- col	Con curr ency	Message Size	MyNETBen ch Measured Latency (ms)	ping Measure d Latency (ms)	Theoretical Latency (ms)	MyNETB ench Efficiency (%)	iperf Efficien cy (%)
TCP	1	1B	0.110705	0.55	0.63	17.57	87.3
TCP	2	1B	0.058607	0.35	0.36	16.27	97.22
TCP	4	1B	0.031419	0.12	0.21	14.96	57.14
TCP	8	1B	0.019337	0.155	0.174	11.11	89.08
UDP	1	1B	0.1092	0.51	0.821	13.30	62.11
UDP	2	1B	0.0554	0.32	0.43	12.88	74.41
UDP	4	1B	0.0308	0.21	0.287	10.73	73.17
UDP	8	1B	0.02152	0.11	0.154	13.97	71.42

IPERF Benchmarking for TCP and UDP:

The following screenshot shows the performance metrics taken from IPerf benchmarking ran on the system.

The speeds are measured in terms of Kilobytes/sec. The first image shows the bandwidth/ throughput for UDP and the second image shows the throughput for TCP protocol.

```
[ 3] 0.0-10.0 sec 672 KBytes 550 Kbits/sec
schoudhari@compute-6:~$ iperf -c 172.16.1.65 -l 1b -b 10000m -u
WARNING: option -l has implied compatibility mode
WARNING: option -b implies udp testing
-----
Client connecting to 172.16.1.65, UDP port 5001
Sending 1 byte datagrams
UDP buffer size: 208 KByte (default)
-----
[ 3] local 172.16.1.6 port 45539 connected with 172.16.1.65 port 5001
[ ID] Interval      Transfer    Bandwidth
[ 3] 0.0-10.0 sec 671 KBytes 549 Kbits/sec
[ 3] Sent 686757 datagrams
schoudhari@compute-6:~$ iperf -c 172.16.1.65 -l 1kb -b 10000m -u
WARNING: option -b implies udp testing
-----
Client connecting to 172.16.1.65, UDP port 5001
Sending 1000 byte datagrams
UDP buffer size: 208 KByte (default)
-----
[ 3] local 172.16.1.6 port 35023 connected with 172.16.1.65 port 5001
[ ID] Interval      Transfer    Bandwidth
[ 3] 0.0-10.0 sec 656 MBytes 550 Mbits/sec
[ 3] Sent 687701 datagrams
[ 3] Server Report:
[ 3] 0.0-10.0 sec 636 MBytes 534 Mbits/sec 0.005 ms 20400/687700 (3%)
[ 3] 0.0-10.0 sec 1 datagrams received out-of-order
schoudhari@compute-6:~$ iperf -c 172.16.1.65 -l 32kb -b 10000m -u
WARNING: option -b implies udp testing
-----
Client connecting to 172.16.1.65, UDP port 5001
Sending 32000 byte datagrams
UDP buffer size: 208 KByte (default)
-----
[ 3] local 172.16.1.6 port 35391 connected with 172.16.1.65 port 5001
[ ID] Interval      Transfer    Bandwidth
[ 3] 0.0-10.0 sec 8.79 GBytes 7.55 Gbits/sec
[ 3] Sent 294822 datagrams
[ 3] Server Report:
[ 3] 0.0-10.0 sec 41.5 MBytes 34.8 Mbits/sec 0.050 ms 265221/294819 (90%)
[ 3] 0.0-10.0 sec 4 datagrams received out-of-order
```

```

schoudhari@compute-6:~$ iperf -c 172.16.1.65 -l 1kb
-----
Client connecting to 172.16.1.65, TCP port 5001
TCP window size: 45.0 KByte (default)
-----
[ 3] local 172.16.1.6 port 53148 connected with 172.16.1.65 port 5001
[ ID] Interval      Transfer    Bandwidth
[ 3] 0.0-10.0 sec  657 MBytes  551 Mbits/sec
schoudhari@compute-6:~$ iperf -c 172.16.1.65 -l 32kb
-----
Client connecting to 172.16.1.65, TCP port 5001
TCP window size: 45.0 KByte (default)
-----
[ 3] local 172.16.1.6 port 53150 connected with 172.16.1.65 port 5001
[ ID] Interval      Transfer    Bandwidth
[ 3] 0.0-10.0 sec  5.86 GBytes  5.03 Gbits/sec
schoudhari@compute-6:~$ iperf -c 172.16.1.65 -l 1b
WARNING: option -l has implied compatibility mode
-----
Client connecting to 172.16.1.65, TCP port 5001
TCP window size: 45.0 KByte (default)
-----
[ 3] local 172.16.1.6 port 53152 connected with 172.16.1.65 port 5001
[ ID] Interval      Transfer    Bandwidth
[ 3] 0.0-10.0 sec   672 KBytes  550 Kbits/sec

```

Conclusion:

Performed benchmarking on different computer parts including CPU, RAM(memory), Disk and Network. Implemented all the benchmark using strong scaling experiment written in C language and evaluated the efficiency, by comparing the output with the values obtained from the benchmark tool and calculated theoretical values.

References:

<http://man7.org/linux/man-pages/man2/connect.2.html>
<http://www2.cs.uidaho.edu/~krings/CS270/Notes.S10/270-F10-28.pdf>
<https://www.programiz.com/c-programming/c-dynamic-memory-allocation>
<https://stackoverflow.com/questions/15731924/udp-multi-client-server-basics>